

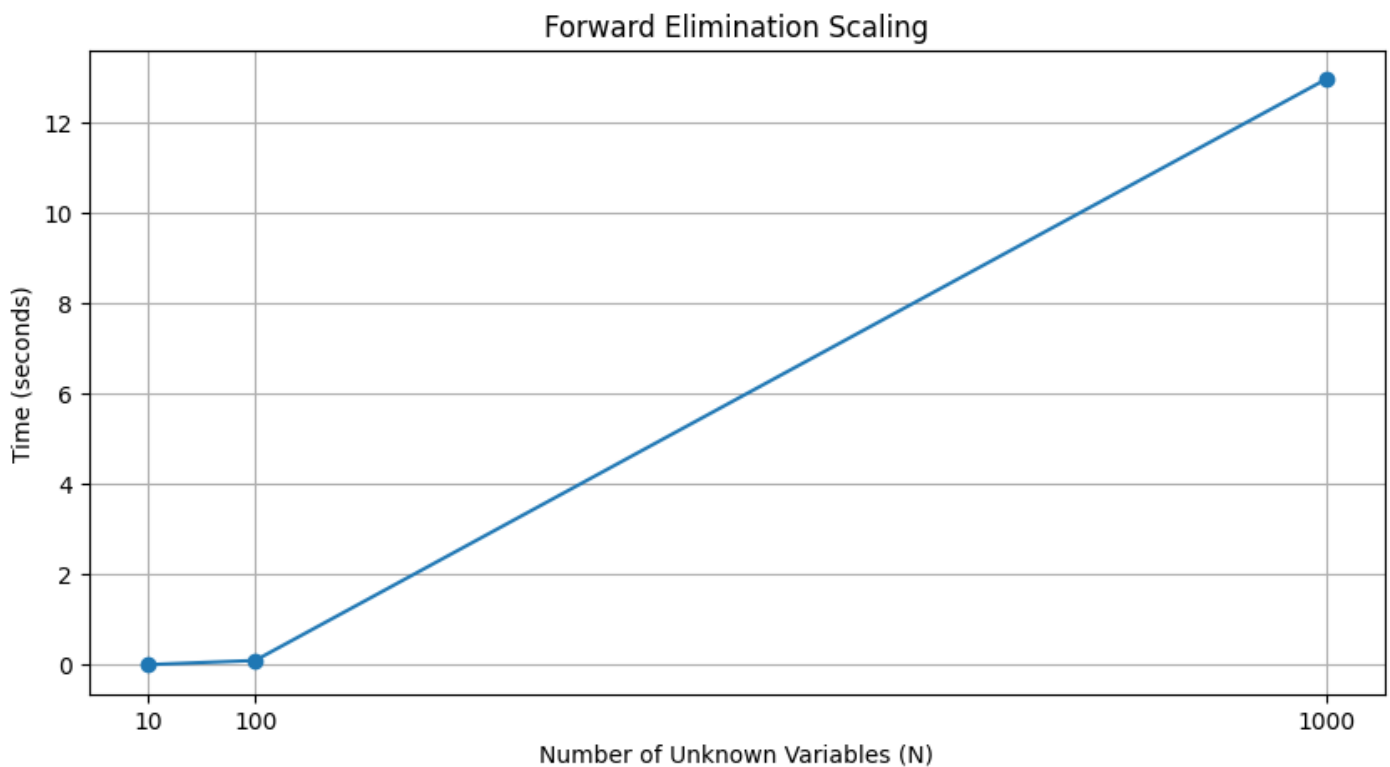
Gaussian Elimination in Parallel

Overview

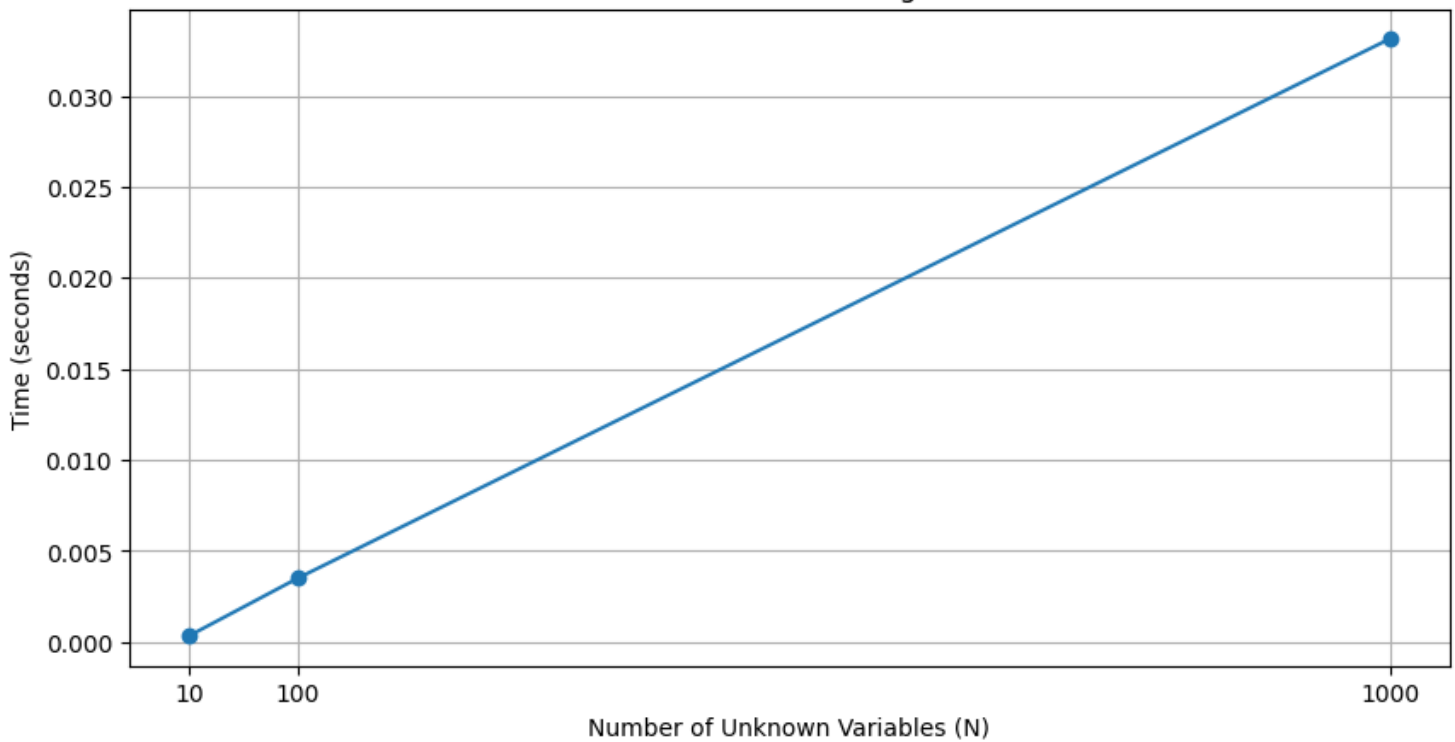
- This program demonstrates that certain tasks do not translate well into parallel.
- We developed a C program that will do gaussian elimination in parallel
- The forward elimination and the backsolving are implemented and measured seperately.
- We run the function on input sizes of random matrices of sizes 10, 100, 1000 to see how the program scales with input.

Results

- The results show that the program does not scale well when it comes to the forward elimination but it doesn't suffer as such with the backsolving.
- This is likely because the forward elimination requires thread synchronization and communication whereas the backsolving does not and is much less complex.



Backsolve Scaling



Code

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <time.h>
#include <math.h>
// Macro for calculating the min value
#define MIN(a, b) ((a) < (b) ? (a) : (b))

// This will help us pass all the required data to the forward_elimination
// function which will be run by each thread
typedef struct {
    float **matrix; // 2D array to store the matrix
    int num_rows; // number of rows
    int num_cols; // number of columns

    int thread_id; // Current thread id to calculate which rows to process
    // total number of threads we are using.
    int num_threads; // Used with thread_id to divide the work

    // mutex to protect the sync_counter
    pthread_mutex_t *mutex;
    pthread_cond_t *cond; // condition used to broadcast to other threads
    // that we are ready to go to the next column

    // keeps track of how many threads reached the critical instructions
    int *sync_counter;
} thread_data_t;
```

```

// This serves a similar function as thread_data_t but for the backsolving
typedef struct {
    float **matrix;
    int num_rows;
    int num_cols;
    float *solution; // stores the solution array.
    int variable_index; // the index that the thread will work on
} backsolve_thread_data_t;

```

```

// Because we don't want to nest another for loop

```

```

void update_row(
    float **matrix,
    int row,
    int start_col,
    int num_cols,
    float factor,
    float *pivot_row
)
{
    int j;
    for (j = start_col; j < num_cols; j++) {
        matrix[row][j] -= factor * pivot_row[j];
    }
}

```

```

// Configured to be used by pthreads

```

```

// each thread will perform forward elimination on it's own section.

```

```

void *forward_elimination(void *arg) {

    // convert the argument back to struct thread_data_t
    thread_data_t *data = (thread_data_t *)arg;

    float factor;
    for (int k = 0; k < data->num_cols - 1; k++) {
        if (k % data->num_threads == data->thread_id) {
            for (int i = k + 1; i < data->num_rows; i++) {
                factor =
                    data->matrix[i][k] / data->matrix[k][k];

                update_row(
                    data->matrix,
                    i, k,
                    data->num_cols,
                    factor,
                    data->matrix[k]
                );
            }
        }

        // lock the mutex before incrementing thread counter
        pthread_mutex_lock(data->mutex);
    }
}

```

```

        (*data->sync_counter)++; // increment the thread counter

        // if not all threads finished on this column
        if (*data->sync_counter < data->num_threads) {
            // wait for them to be done then proceed to next column
            pthread_cond_wait(data->cond, data->mutex);
        } else { // else

            // start over and proceed to next column
            *data->sync_counter = 0;

            // let all the other threads know we are done
            pthread_cond_broadcast(data->cond);
        }

        // unlock the mutex
        pthread_mutex_unlock(data->mutex);
    }

    pthread_exit(NULL);
}

// Configured to be used by pthreads
// This will be run by each thread in the backsolve function
void *backsolve_thread(void *arg) {
    backsolve_thread_data_t *data = (backsolve_thread_data_t *)arg;
    int i = data->variable_index;
    float sum = 0;
    for (int j = i + 1; j < data->num_cols - 1; j++) {
        sum += data->matrix[i][j] * data->solution[j];
    }
    data->solution[i] = (data->matrix[i][data->num_cols - 1] - sum) / data->matrix[i][i];
    pthread_exit(NULL);
}

// Function to do the backsolving
float *backsolve(float **matrix, int num_rows, int num_cols) {

    // initialize the array to store the solution
    float *solution = malloc(num_rows * sizeof(float));

    // initialize an array of threads equal to the number of unknowns
    pthread_t *threads = malloc(num_rows * sizeof(pthread_t));

    // Array of the arguments for each thread function call
    backsolve_thread_data_t *thread_data =
        malloc(num_rows * sizeof(backsolve_thread_data_t));

    // configure the arguments for each thread and create them
    for (int thread_id = num_rows - 1; thread_id >= 0; thread_id--) {
        thread_data[thread_id].matrix = matrix;
        thread_data[thread_id].num_rows = num_rows;
        thread_data[thread_id].num_cols = num_cols;
    }
}

```

```

        thread_data[thread_id].solution = solution;
        thread_data[thread_id].variable_index = thread_id;

        pthread_create(
            &threads[thread_id],
            NULL,
            backsolve_thread,
            &thread_data[thread_id]
        );
    }

    // Wait for all the threads to complete before ending and returning
    // the solution
    for (int thread_id = num_rows - 1; thread_id >= 0; thread_id--) {
        pthread_join(threads[thread_id], NULL);
    }

    // Free the allocated memory
    free(threads);
    free(thread_data);

    return solution;
}

// Create a num_rows by num_cols matrix with random floats or predefined matrix
float **create_matrix(int num_rows, int num_cols, int predefined)
{
    float **matrix = malloc(num_rows * sizeof(float *));
    if (predefined) {
        // predefined matrix that we know the answer to for testing
        float predefined_matrix[5][6] = {
            {2, -1, 1, 3, 1, 10},
            {1, 3, 2, -1, 2, 5},
            {1, -1, 2, 4, 1, 7},
            {3, 2, -1, 1, 2, 12},
            {2, 1, 3, -2, 1, 3}
        };

        for (int row = 0; row < num_rows; row++) {
            matrix[row] = malloc(num_cols * sizeof(float));
            for (int col = 0; col < num_cols; col++) {
                matrix[row][col] = predefined_matrix[row][col];
            }
        }
    } else {
        for (int row = 0; row < num_rows; row++) {
            matrix[row] = malloc(num_cols * sizeof(float));
            for (int col = 0; col < num_cols; col++) {
                matrix[row][col] = (float)rand() / RAND_MAX * 2000 - 1000;
            }
        }
    }
    return matrix;
}

```

```

// Frees all the memory created for a matrix
void free_matrix(float **matrix, int num_rows)
{
    // free inner rows of the matrix
    int row;
    for (row = 0; row < num_rows; row++) {
        free(matrix[row]);
    }

    // Free the outter memory of the matrix
    free(matrix);
}

// Function to verify the solution
void verify_solution(float **matrix, float *solution, int num_rows, int num_cols)
{
    float sum;
    for (int i = 0; i < num_rows; i++) {
        sum = 0;
        for (int j = 0; j < num_cols - 1; j++) {
            sum += matrix[i][j] * solution[j];
        }
        if (fabs(sum - matrix[i][num_cols - 1]) > 1e-5) {
            printf("Solution is incorrect!\n");
            return;
        }
    }
    printf("Solution is correct!\n");
}

int main()
{
    // initialize our thread sync variables
    pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
    pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
    int sync_counter = 0;

    // The input matrix sizes to test
    int input_matrix_sizes[] = { 10, 100, 1000 };

    for (int size_index = 0; size_index < 3; size_index++) {
        int num_rows = input_matrix_sizes[size_index];
        int num_cols = num_rows + 1;

        float **matrix = create_matrix(num_rows, num_cols, 0);

        // We store the threads and their respective inputs in arrays
        // So that we can iterate through them and execute them.
        pthread_t *threads = malloc(num_rows * sizeof(pthread_t));
        thread_data_t *thread_data =
            malloc(num_rows * sizeof(thread_data_t));
    }
}

```

```

clock_t start, end; // setup our timer
double cpu_time_used;

start = clock(); // record start time of forward elimination

for (int thread_id = 0; thread_id < num_rows; thread_id++) {

    // Initialize all the inputs for this thread
    thread_data[thread_id].matrix = matrix;
    thread_data[thread_id].num_rows = num_rows;
    thread_data[thread_id].num_cols = num_cols;
    thread_data[thread_id].thread_id = thread_id;
    thread_data[thread_id].num_threads = num_rows;
    thread_data[thread_id].mutex = &mutex;
    thread_data[thread_id].cond = &cond;
    thread_data[thread_id].sync_counter = &sync_counter;

    // start the thread and pass the appropriate inputs
    // forward_elimination() is the function the threads run
    pthread_create(
        &threads[thread_id],
        NULL,
        forward_elimination,
        &thread_data[thread_id]
    );
}

// Wait for all threads to complete before proceeding
for (int thread_id = 0; thread_id < num_rows; thread_id++) {
    pthread_join(threads[thread_id], NULL);
}

end = clock(); // record end time of forward elimination

// calculate time elapsed and convert to CPU time
cpu_time_used = ((double)(end - start)) / CLOCKS_PER_SEC;

// print results for forward elimination
printf(
    "Time for forward elimination N=%d: %f seconds\n",
    num_rows,
    cpu_time_used
);

start = clock(); // record start time of backsolving
backsolve(matrix, num_rows, num_cols);
end = clock(); // record end time of backsolving

// calculate time elapsed and convert to CPU time
cpu_time_used = ((double)(end - start)) / CLOCKS_PER_SEC;

```

```

// print results for backsolving
printf(
    "Time taken for backsolve with N=%d: %f seconds\n",
    num_rows,
    cpu_time_used
);

// Used to ensure the gaussian elimination is correct
// Capture the solution array
// float* solution = backsolve(matrix, num_rows, num_cols);
// Print the solution
// printf("Solution:\n");
// for (int i = 0; i < num_rows; i++) {
//     printf("%f ", solution[i]);
// }
// printf("\n");

// Verify the solution
// verify_solution(matrix, solution, num_rows, num_cols);

// Free the solution array
// free(solution);

// No memory left behind.
free_matrix(matrix, num_rows);
free(threads);
free(thread_data);
}

return 0;
}

```

Raw outputs

```

ziadarafat@dogearchvm ~/D/s/parallel-programming (main)> cd Gaussian_Elimination/
ziadarafat@dogearchvm ~/D/s/p/Gaussian_Elimination (main)> make
gcc -O0 -Wall -Werror -Wextra -pedantic -pthread -lpthread -o gaussian_elimination gauss
ziadarafat@dogearchvm ~/D/s/p/Gaussian_Elimination (main)> ./gaussian_elimination
Time for forward elimination N=10: 0.001633 seconds
Time taken for backsolve with N=10: 0.000352 seconds
Time for forward elimination N=100: 0.095059 seconds
Time taken for backsolve with N=100: 0.003582 seconds
Time for forward elimination N=1000: 12.931657 seconds
Time taken for backsolve with N=1000: 0.034218 seconds
ziadarafat@dogearchvm ~/D/s/p/Gaussian_Elimination (main)> ./gaussian_elimination
Time for forward elimination N=10: 0.001544 seconds

```



```
Time taken for backsolve with N=10: 0.000303 seconds
Time for forward elimination N=100: 0.090941 seconds
Time taken for backsolve with N=100: 0.003233 seconds
Time for forward elimination N=1000: 13.024221 seconds
Time taken for backsolve with N=1000: 0.032342 seconds
ziadarafat@dogearchvm ~/D/s/p/Gaussian_Elimination (main)> ./gaussian_elimination
Time for forward elimination N=10: 0.001593 seconds
Time taken for backsolve with N=10: 0.000289 seconds
Time for forward elimination N=100: 0.087520 seconds
Time taken for backsolve with N=100: 0.003795 seconds
Time for forward elimination N=1000: 12.957925 seconds
Time taken for backsolve with N=1000: 0.032653 seconds
ziadarafat@dogearchvm ~/D/s/p/Gaussian_Elimination (main)> ./gaussian_elimination
Time for forward elimination N=10: 0.001633 seconds
Time taken for backsolve with N=10: 0.000289 seconds
Time for forward elimination N=100: 0.092816 seconds
Time taken for backsolve with N=100: 0.003512 seconds
Time for forward elimination N=1000: 12.753179 seconds
Time taken for backsolve with N=1000: 0.032123 seconds
ziadarafat@dogearchvm ~/D/s/p/Gaussian_Elimination (main)> ./gaussian_elimination
Time for forward elimination N=10: 0.001430 seconds
Time taken for backsolve with N=10: 0.000427 seconds
Time for forward elimination N=100: 0.099994 seconds
Time taken for backsolve with N=100: 0.003431 seconds
Time for forward elimination N=1000: 13.095672 seconds
Time taken for backsolve with N=1000: 0.034295 seconds
```